

John Doe

+1 (555) 123-4567 | San Francisco, CA | john.doe@example.com | github.com/johndoe | linkedin.com/in/johndoe | johndoe.dev

Summary

Software engineer with 4+ years of experience building and deploying cloud-native applications and machine learning pipelines. Skilled in designing scalable distributed systems on AWS and GCP, fine-tuning large language models, and delivering production-grade APIs. Strong background in MLOps, infrastructure-as-code, and cross-functional collaboration.

Skills

- Languages:** Python, TypeScript, Go, SQL, Bash, YAML
- AI & ML:** PyTorch, TensorFlow, Hugging Face Transformers, LangChain, OpenAI API, scikit-learn, Pandas, NumPy
- Cloud & DevOps:** AWS (EC2, S3, Lambda, SageMaker), GCP (Vertex AI, BigQuery, Cloud Run), Docker, Kubernetes, Terraform, CI/CD, GitHub Actions

Experience

Senior Machine Learning Engineer Jun 2023 — Present
Cloudware AI San Francisco, CA

- Designed and deployed a RAG-based document intelligence platform on AWS using LangChain, FAISS, and GPT-4, reducing manual document review time by 70%
- Built and maintained MLOps pipelines on SageMaker for model training, evaluation, and A/B deployment across 12 production models serving 5M+ daily requests
- Led migration of monolithic inference service to containerized microservices on EKS, improving p99 latency by 40% and cutting infrastructure costs by 30%
- Mentored 3 junior engineers on LLM fine-tuning workflows and cloud cost optimization best practices

Machine Learning Engineer Aug 2021 — May 2023
DataBridge Corp Austin, TX

- Fine-tuned open-source LLMs (LLaMA 2, Mistral) using LoRA and PEFT on GCP Vertex AI for domain-specific NLP tasks, achieving 18% improvement over baseline
- Developed real-time anomaly detection system using streaming data from Pub/Sub and BigQuery ML, reducing incident response time by 55%
- Automated model retraining and monitoring pipelines with Airflow and Terraform, eliminating 15+ hours of manual ops work per week
- Collaborated with product and data teams to define ML feature roadmap and translate business requirements into model specifications

Cloud Infrastructure Engineer May 2020 — Jul 2021
NexGen Systems Seattle, WA

- Provisioned and managed multi-region AWS infrastructure using Terraform and CloudFormation for a SaaS platform serving 200K users
- Implemented CI/CD pipelines with GitHub Actions and AWS CodePipeline, reducing deployment cycle from 2 days to under 2 hours
- Optimized cloud spend by 35% through reserved instance planning, S3 lifecycle policies, and right-sizing of EC2 workloads

Projects

Creator & Maintainer, OpenInfer (openinfer.dev) Jan 2024 — Present

- Open-source LLM inference server built with Python and FastAPI, supporting OpenAI-compatible endpoints for self-hosted models with 1.2K+ GitHub stars
- Implemented dynamic batching, quantization (GPTQ, AWQ), and multi-GPU support, achieving 3x throughput vs naive serving baseline
- Deployed reference infrastructure on AWS with Terraform and Docker; full observability via Prometheus and Grafana dashboards

Lead Developer, CloudCost Analyzer Sep 2023 — Dec 2023

- Built an AI-powered AWS cost analysis tool using GPT-4 and the AWS Cost Explorer API to surface savings recommendations in plain English
- Reduced average cloud bill by 22% across 10 pilot teams by automating idle resource detection and right-sizing suggestions

Education

University of California, San Diego

Master of Science, Computer Science — GPA: 4.0

San Diego, CA

Sep 2020 — Jun 2022

University of California, San Diego

Bachelor of Science, Computer Science

San Diego, CA

Sep 2016 — Jun 2020

Certifications

- AWS Certified Machine Learning – Specialty
- Google Professional Machine Learning Engineer
- AWS Certified Solutions Architect – Associate
- Certified Kubernetes Administrator (CKA)